

July 29, 2026

RE: Member of Technical Staff – Sovereign AI, Cohere

Dear Cohere Hiring Team,

I am excited to apply for the Member of Technical Staff – Sovereign AI role advertised through Cohere’s careers site. I am a Toronto-based staff software developer and enterprise architect who works across applied LLM systems, secure platform engineering, and AI governance. Cohere’s research-and-engineering model is especially compelling to me: I do my best work when I can test novel ideas rigorously, turn the useful ones into dependable software, and deploy them for consequential real-world use.

At RBC, I am building ASB Assist, an agentic LLM + RAG system that turns natural-language requirements into compliant architecture diagrams and enterprise service recommendations. My work spans Python orchestration, retrieval, tool integration, context design, offline evaluation, grounding and citation checks, guardrails, and end-to-end auditability. It builds on ASB Builder, the deterministic platform I designed and shipped from zero to production using TypeScript, React, and REST APIs. I authored nearly all of its production code and the complete recommendation engine, which maps more than 1,000 architecture patterns to real-time component selections. Building both systems has taught me to combine model capabilities with deterministic controls where security and correctness are non-negotiable.

Sovereignty is also central to my current architecture work. I operationalized AI governance spanning model risk, Responsible AI, IAM, data residency, regulatory requirements, and auditable deployment across AWS, GCP, and air-gapped OpenShift environments. Separately, I architected an OpenTelemetry data platform that ingests 150 TB per day across more than 10,000 systems using Vector, Fluent Bit, Kafka, and Logstash. That work required distributed-systems judgment around throughput, backpressure, failure recovery, observability, and security-constrained operations—the same production discipline needed to serve mission-critical AI.

I bring hands-on ML experimentation as well. I use Python, PyTorch, Hugging Face, FSDP/DDP, DeepSpeed ZeRO, mixed precision, gradient checkpointing, experiment tracking, and reproducible evaluation in distributed LLM training prototypes. Earlier, as HealthCard’s first engineer, I deployed transformer document classification behind monitored inference APIs and built the surrounding cloud platform. My graduate work included research on call-stack decision algorithms and teaching algorithms and advanced cloud computing; I also published peer-reviewed work on privacy-preserving decentralized storage. At RBC, I now lead five engineers through research spikes, implementation, review, testing, and operational handoff while remaining deeply involved in the code.

I would welcome the opportunity to help Cohere research, build, and scale agentic systems that keep sensitive Canadian workloads secure and deployable under sovereign constraints. The combination of frontier-model access, public-interest applications, and full-stack ownership makes this a rare fit for how I want to contribute.

Thank you for your consideration. I would be glad to discuss how my applied AI, distributed-systems, and secure-platform experience could support Cohere’s Sovereign AI team.

Best regards,
Dhruv Doshi