

May 15, 2026

RE: Senior Engineer, Applied AI – FICO

Dear Hiring Team,

I am applying for the Senior Engineer role on FICO's Applied AI team. The mandate—production agents, RAG, rigorous evaluation, and responsible AI on a decision-management platform—matches what I have been building at RBC: LLM-powered workflows, enterprise guardrails, and observability at scale in a regulated financial environment. FICO's focus on ethical, high-impact AI in fraud, decision automation, and customer outcomes is exactly where I want to do my next chapter of work.

Today I am building ASB Assist, a conversational agentic system on top of our enterprise blueprint platform. Users describe requirements in natural language; the system generates compliant architecture outputs via an LLM + RAG pipeline with LangChain orchestration, MCP-based tool integration, and multi-step agent workflows. I own prompt and context engineering, tool design, offline quality iteration, and alignment with security patterns (OAuth, Zero-Trust, agent gateways). In parallel I defined RBC's enterprise AI governance framework—Responsible AI, model risk, audit trails, OSFI/GDPR constraints—so production AI ships with guardrails, not afterthoughts. I also architected institution-wide observability (150 TB/day, OpenTelemetry, Kafka) that standardised how 20+ teams monitor distributed services, including AI/ML operational use cases. Underneath ASB Assist I delivered ASB Builder (patent-pending), a rules-based platform I built 0-to-1 in TypeScript/React that now serves engineers org-wide.

Earlier I deployed transformer-based document classification to production on AWS EKS at HealthCard—inference APIs, monitoring, and measurable reduction in manual review—which taught me how to close the loop from model behaviour to business metrics. My MASc (GPA 3.9, Dalhousie) and graduate TA work in algorithms and cloud computing ground the engineering rigour I bring to eval design and systems trade-offs (latency, cost, reliability).

I would welcome the opportunity to help FICO ship agentic and LLM capabilities with the evaluation discipline and responsible AI standards your VP describes. Thank you for your consideration; I am happy to discuss specific examples of agent design, RAG architecture, or governance patterns at your convenience.

Best regards,



Dhruv Doshi